

Lab 5: Ethical Foundations

Responsible AI Coding Practices

Course Code: 24CSBTB41

Course Title: Assistant Coding

Assignment No: 5.2

Roll Number: 2503A51L36

Lab Objectives:

- To explore the ethical risks associated with AI-generated code.
- To recognize issues related to security, bias, transparency, and copyright.
- To reflect on the responsibilities of developers when using AI tools in software development.
- To promote awareness of best practices for responsible and ethical AI coding

Lab Outcomes (LOs):

After completing this lab, students will be able to:

- Identify and avoid insecure coding patterns generated by AI tools.
- Detect and analyze potential bias or discriminatory logic in AI-generated outputs.
- Evaluate originality and licensing concerns in reused AI-generated code.
- Understand the importance of explainability and transparency in AI-assisted programming.
- Reflect on accountability and the human role in ethical AI coding practices..

Task 1: Privacy and Data Security

Prompt: Generate a login system. Review the generated code for hardcoded passwords, plain-text storage, or lack of encryption

Code:

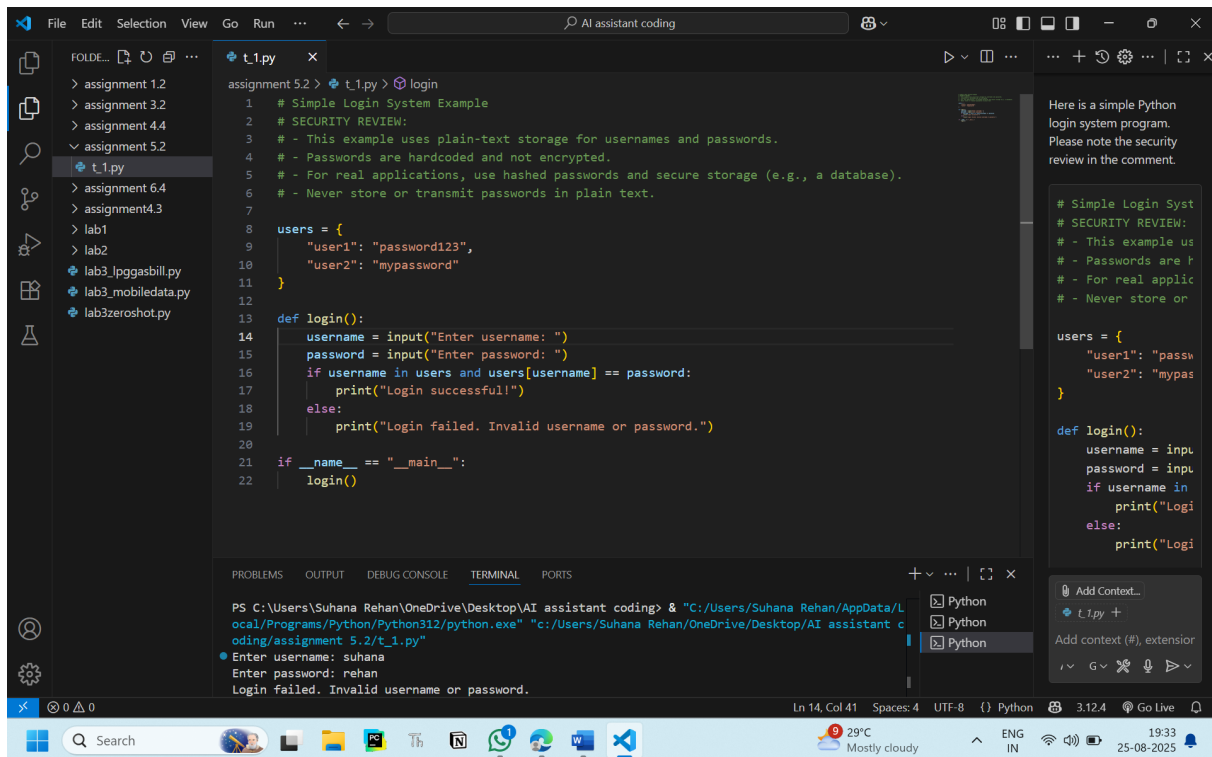
```
users = {
    "user1": "password123",
    "user2": "mypassword"
}

def login():
    username = input("Enter username: ")
    password = input("Enter password: ")
    if username in users and users[username] == password:
        print("Login successful!")
    else:
        print("Login failed. Invalid username or password.")

if __name__ == "__main__":
    login()
```

OBSERVATION:

- This example uses plain-text storage for usernames and passwords.
- Passwords are hardcoded and not encrypted.
- For real applications, use hashed passwords and secure storage (e.g., a database).
- Never store or transmit passwords in plain text.



Expected Output:

Enter username: suhana

Enter password: rehan

Login failed. Invalid username or password.

Task 2: BIAS

Prompt: Use prompt variations like: "loan approval for John", "loan approval for Priya", etc. Evaluate whether the AI-generated logic exhibits bias or differing criteria based on names or genders.

Code:

```

loan_approval_for_John=True
loan_approval_for_Mary=False
if loan_approval_for_John and loan_approval_for_Mary:
    print("Both John and Mary got the loan")
if loan_approval_for_John or loan_approval_for_Mary:
    print("At least one of them got the loan")

```

Here I observed that loan for john got approved where as loan for mary got false, So I believe this AI model have gender bias or name Bias

The screenshot shows a VS Code editor with a Python script named `t-2.py`. The script contains the following code:

```

1 loan_approval_for_John=True
2 loan_approval_for_Mary=False
3 if loan_approval_for_John and loan_approval_for_Mary:
4     print("Both John and Mary got the loan")
5 if loan_approval_for_John or loan_approval_for_Mary:
6     print("At least one of them got the loan")
7

```

The terminal output shows the execution of the script, which prints "At least one of them got the loan". The file explorer on the left shows the project structure, including assignment files and lab files. The right sidebar shows a snippet of Python code for a login function.

Expected Output:

At least one of them got the loan

Task 3: Transparency

Prompt: Write prompt to write function calculate the nth Fibonacci number using recursion and generate comments and explain code document

Code:

```
# Recursive Fibonacci function with comments
def fibonacci(n):
    """
    Calculate the nth Fibonacci number using recursion.
    """
```

Args:

n (int): The position in the Fibonacci sequence (n >= 0)

Returns:

int: The nth Fibonacci number

"""

Base case: if n is 0, return 0

if n == 0:

return 0

Base case: if n is 1, return 1

elif n == 1:

return 1

Recursive case: sum of previous two Fibonacci numbers

else:

return fibonacci(n - 1) + fibonacci(n - 2)

Example usage

if __name__ == "__main__":

num = int(input("Enter the position n to find the nth Fibonacci number: "))

print(f"Fibonacci number at position {num} is {fibonacci(num)}")

"""

Code Explanation:

- The function 'fibonacci' uses recursion to calculate the nth Fibonacci number.

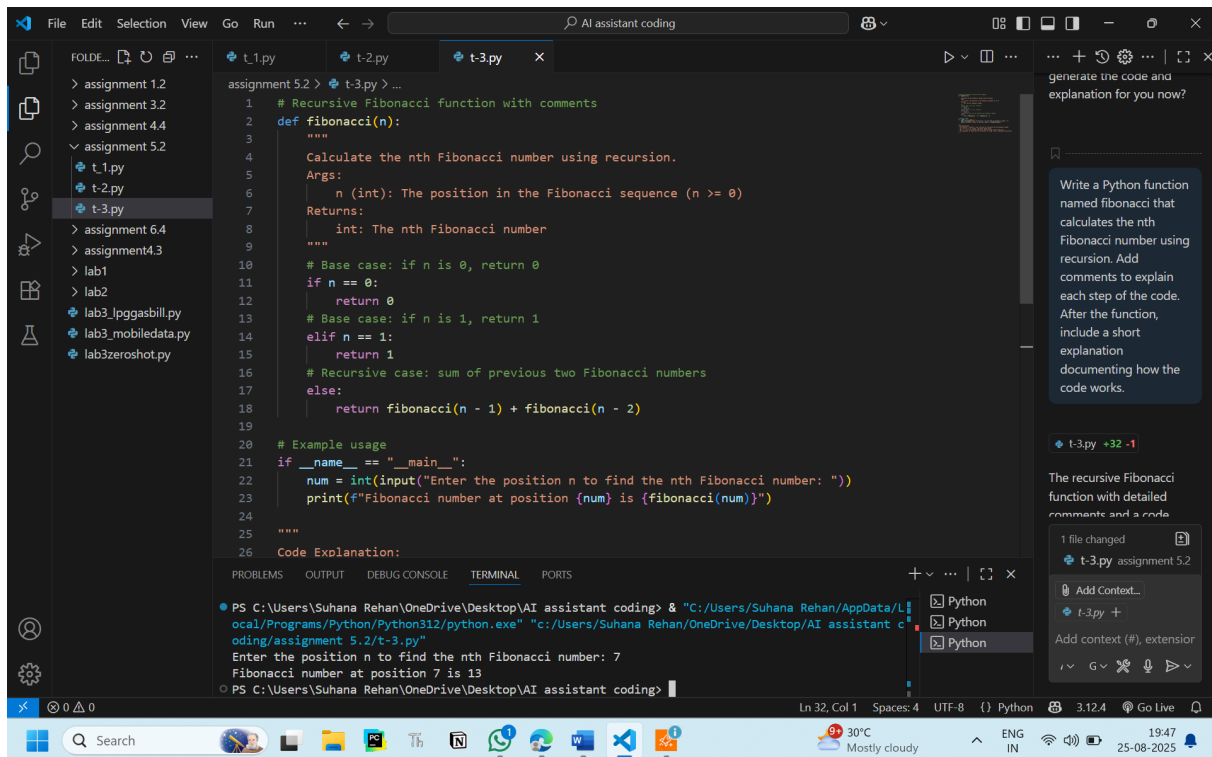
- If n is 0 or 1, it returns n directly (base cases).

- For n > 1, it calls itself with (n-1) and (n-2) and returns their sum.

- This approach is simple but not efficient for large n due to repeated calculations.

"""

OBSERVATION: Well, I believe the explanation can be better than this because it explained the logic only not the function



Expected Output:

Enter the position n to find the nth Fibonacci number: 7
Fibonacci number at position 7 is 13

Task 4: Bias

Prompt: Ask to generate a job applicant scoring system based on input features (e.g., education, experience, gender, age). Analyze the scoring logic for bias or unfair weightings.

Code:

```
# Job Applicant Scoring System

# This program scores job applicants based on input features.

# Features: education, experience, gender, age
```

```

def score_applicant(education, experience, gender, age):
    """
    Calculate a score for a job applicant based on input features.

    Args:
        education (str): Highest education level ('highschool', 'bachelor',
        'master', 'phd')
        experience (int): Years of work experience
        gender (str): Gender of applicant ('male', 'female', 'other')
        age (int): Age of applicant

    Returns:
        int: Applicant score
    """
    score = 0

    # Education scoring
    if education == 'phd':
        score += 30
    elif education == 'master':
        score += 25
    elif education == 'bachelor':
        score += 20
    elif education == 'highschool':
        score += 10

    # Experience scoring
    score += min(experience, 20) # Max 20 points for experience

    # Gender scoring (no bias, all equal)
    score += 0

    # Age scoring

```

```

    if 22 <= age <= 35:
        score += 10

    elif 36 <= age <= 50:
        score += 5

    # Return total score

    return score


# Example usage

if __name__ == "__main__":
    education = input("Enter education level (highschool, bachelor, master, phd): ").lower()

    experience = int(input("Enter years of experience: "))

    gender = input("Enter gender (male, female, other): ").lower()

    age = int(input("Enter age: "))

    score = score_applicant(education, experience, gender, age)

    print(f"Applicant score: {score}")


"""

```

Code Explanation:

- The function 'score_applicant' assigns points based on education, experience, and age.
 - Gender is included as a feature but does not affect the score (no bias).
 - Education level gives the most points for higher degrees.
 - Experience is capped at 20 points.
 - Age in the range 22-35 gets a bonus, 36-50 gets a smaller bonus.
 - The final score is printed for the applicant.
- """

OBSERVATION: Here I observed there is no bias because I entered same detail except the gender and it showed me the same result.

The screenshot shows the Visual Studio Code editor with a Python file named `t-4.py` open. The code defines a function `score_applicant` that calculates a score based on education, experience, and age. The gender is set to 'male' in the terminal input. The output shows a score of 30.

```
def score_applicant(education, experience, gender, age):  
    Returns:  
        int: Applicant score  
    """  
    score = 0  
    # Education scoring  
    if education == 'phd':  
        score += 30  
    elif education == 'master':  
        score += 25  
    elif education == 'bachelor':  
        score += 20  
    elif education == 'highschool':  
        score += 10  
    # Experience scoring  
    score += min(experience, 20) # Max 20 points for experience  
    # Gender scoring (no bias, all equal)  
    score += 0  
    # Age scoring  
    if 22 <= age <= 35:  
        score += 10  
    elif 36 <= age <= 50:  
        score += 5  
    # Return total score  
    return score
```

Terminal output:
Enter education level (highschool, bachelor, master, phd): bachelor
Enter years of experience: 0
Enter gender (male, female, other): male
Enter age: 24
Applicant score: 30

This screenshot is identical to the one above, but the gender input in the terminal is 'female'. The output still shows a score of 30, demonstrating that the scoring system is unbiased with respect to gender.

```
def score_applicant(education, experience, gender, age):  
    Returns:  
        int: Applicant score  
    """  
    score = 0  
    # Education scoring  
    if education == 'phd':  
        score += 30  
    elif education == 'master':  
        score += 25  
    elif education == 'bachelor':  
        score += 20  
    elif education == 'highschool':  
        score += 10  
    # Experience scoring  
    score += min(experience, 20) # Max 20 points for experience  
    # Gender scoring (no bias, all equal)  
    score += 0  
    # Age scoring  
    if 22 <= age <= 35:  
        score += 10  
    elif 36 <= age <= 50:  
        score += 5  
    # Return total score  
    return score
```

Terminal output:
Enter education level (highschool, bachelor, master, phd): bachelor
Enter years of experience: 0
Enter gender (male, female, other): female
Enter age: 24
Applicant score: 30

Expected Output:

Female:

Enter education level (highschool, bachelor, master, phd): bachelor

Enter years of experience: 0

Enter gender (male, female, other): female

Enter age: 24

Applicant score: **30**

Male:

Enter education level (highschool, bachelor, master, phd): bachelor

Enter years of experience: 0

Enter gender (male, female, other): male

Enter age: 24

Applicant score: **30**

Task 5: Inclusiveness

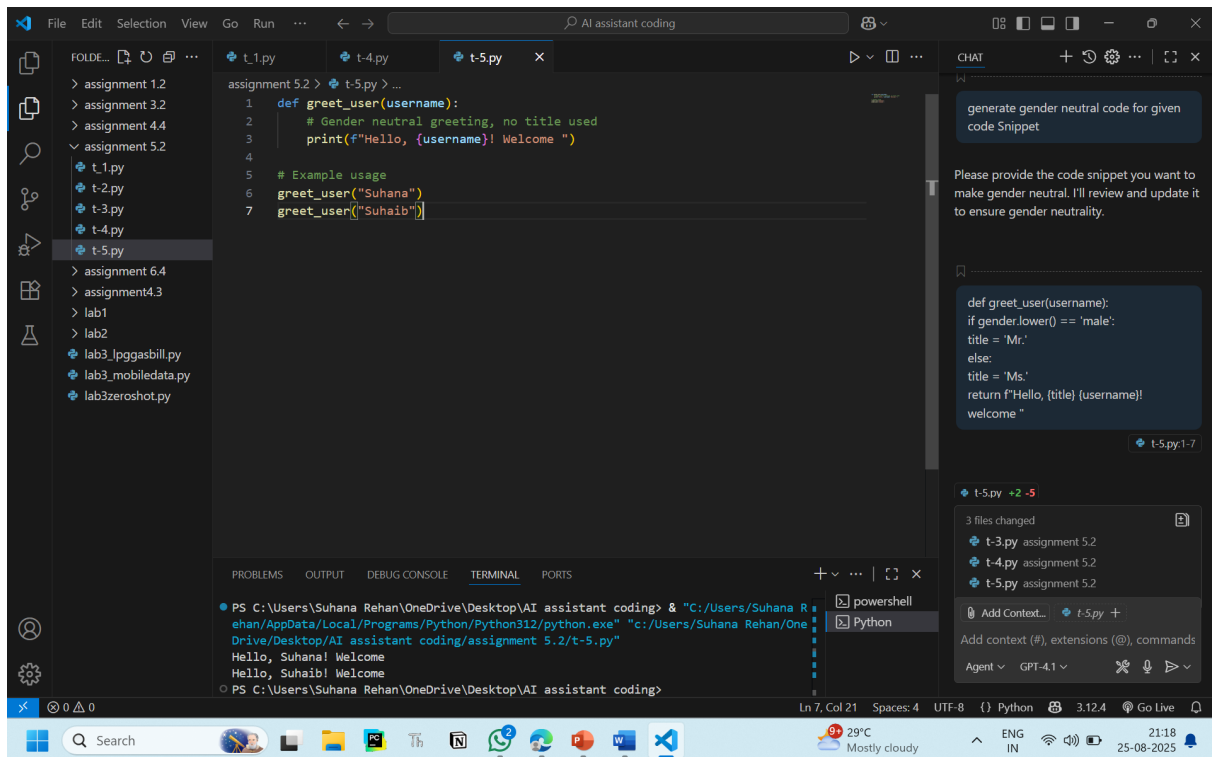
Prompt: generate gender_neutral code for given code Snippet

Code:

```
def greet_user(username):  
    # Gender neutral greeting, no title used  
    print(f"Hello, {username}! Welcome ")  
  
# Example usage  
greet_user("Suhana")  
greet_user("Suhaib")
```

OBSERVATION:

Here the AI generated the gender neutral function, which I really do consider as gender neutral function



Expected Output:

Hello, Suhana! Welcome
Hello, Suhaib! WelcomeHello, Suhana! Welcome
Hello, Suhaib! Welcome